

# STAT-5330 Project: Analyzing Real Estate Market Trends and Price Prediction

Yidan Chen

## **Abstract**

In this project, We present a clustering analysis of New York City's real estate market, focusing on geographic and pricing patterns. Properties are grouped into distinct clusters using k-means, revealing differences in housing characteristics across boroughs.

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Literature review</b>	<b>1</b>
<b>3</b>	<b>Analysis &amp; Results</b>	<b>2</b>
<b>4</b>	<b>Discussion and Conclusion</b>	<b>6</b>
<b>A</b>	<b>Code</b>	<b>7</b>

# 1 Introduction

The housing market is a cornerstone of economic activity and urban development. In New York City, where housing prices vary dramatically, understanding the factors that influence these prices is crucial for buyers, sellers, and policymakers. This study addresses two key questions: **What are the most important factors affecting housing prices?** and **Can we develop an accurate model to predict home prices based on property characteristics?**

To answer these questions, we use a dataset of real estate listings from New York City. The analysis begins with clustering methods, which segment the housing market based on property characteristics and location, highlighting distinct price patterns. We then apply regression analysis to identify significant price determinants, including property size, bedrooms, bathrooms, property type, and location. These methods help us compare multiple predictive models, evaluating their performance using metrics such as AIC and Adjusted  $R^2$ .

## 2 Literature review

Clustering techniques are widely used in housing market studies to identify market segments by grouping properties with similar price trends. [1][2] This study extends these methods by employing location-based clustering to analyze spatial price variations in New York City. Unlike traditional zoning by boroughs (e.g., Manhattan or Brooklyn), we use geographic coordinates to define clusters, providing a finer-grained perspective on pricing patterns.

Regression analysis is a key tool for predicting housing prices. Prior research has demonstrated the importance of attributes like square footage and property type, often using interaction terms to improve accuracy. For instance, interactions between property size and neighborhood characteristics have been shown to enhance pre-

dictions. [3] Our approach combines these predictors with location-based clusters to capture both property-level and spatial effects.

### 3 Analysis & Results

We began by exploring the dataset to understand key variable distributions and relationships. Data cleaning identified anomalies, such as high-priced outliers, duplicate properties with differing prices, and fractional values for bathrooms and bedrooms. To address skewness, we excluded properties above the 95th price percentile. Initially, the price distribution was right-skewed, dominated by lower-priced properties and a few outliers. After removing duplicates and applying a logarithmic transformation to prices, the distribution became more symmetrical, improving its suitability for modeling.




Figure 1: Initial Price Distribution

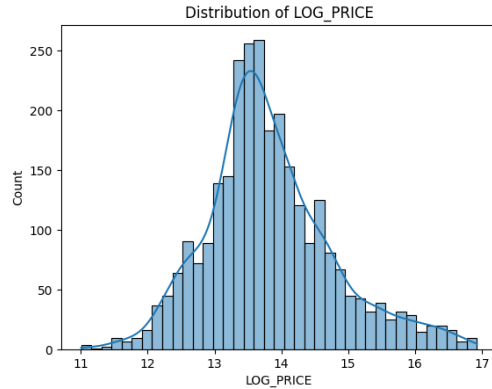





Figure 2: Logarithmic transformation



We then investigated the relationship between price and geographic location (latitude and longitude). Using K-means clustering, we segmented properties into geographic clusters based on their coordinates. This approach provided a more precise spatial segmentation compared to administrative sub-local divisions, revealing distinct price patterns across different areas. High-priced clusters typically aligned with central urban areas, while lower-priced clusters were located on the peripheries. Furthermore, clustering properties based on price and size revealed a clear segmentation of larger, high-value properties versus smaller, more affordable options. These clusters were visualized using scatter plots (price vs. property square footage) and geographic maps. Consistent color coding across clusters facilitated the identification of spatial and numerical patterns.




In the regression analysis, we utilized **R** to construct and refine models predicting property prices. Initially, we compared models incorporating **SUBLOCALITY** (administrative divisions) and **CLUSTER** (derived from K-means clustering of latitude and longitude). The cluster-based model demonstrated a better fit with a lower AIC, emphasizing the effectiveness of spatial clustering in capturing price variations.

Through backward selection, we refined the model to its final form:

$$PRICE \sim PROPERTYSQFT + BEDS + BATH + TYPE + CLUSTER.$$

This model had the lowest AIC and identified key factors influencing prices: First is Square footage (PROPERTYSQFT), which is a highly significant factor, showing a strong positive influence on price. Then they are Bathrooms (BATH) and bedrooms (BEDS), both significantly contributed to price predictions, with bathrooms having a larger impact. For certain clusters, like CLUSTER2 and CLUSTER3, were highly significant, capturing spatial price differences, while others (e.g., CLUSTER1) were not as impactful. Property types (TYPE): Categories like "condo for sale" and "townhouse for sale" significantly influenced prices, whereas others, like "coming soon," had limited impact.



This refined model validated the importance of square footage, property characteristics, and spatial clustering in understanding and predicting housing prices, while identifying factors with limited significance.

## 4 Discussion and Conclusion

This project integrates clustering and regression to evaluate how clustering enhances price prediction accuracy. By exploring various models, we uncovered key factors influencing housing prices in New York City and demonstrated the value of combining these techniques. Square footage (**PROPERTYSQFT**), bathrooms (**BATH**), and bedrooms (**BEDS**) emerged as key predictors, with property type and location-based clusters providing additional insights. Clustering effectively revealed geographic price variations beyond traditional administrative boundaries.

Limitations include assumptions of cluster homogeneity, linearity in regression, and data quality challenges such as outliers. Future work could explore non-linear models and incorporate socio-economic or temporal data to improve predictions.

## References

- [1] T. Skovajsa. (2023). A Review of Clustering Methods for Housing Market Segmentation. *Review of European and Comparative Real Estate*, 11(3), 22-35. Retrieved from <https://sciendo.com/pdf/10.2478/remav-2023-0022>
- [2] S. Hwang and J.-C. Thill. (2007). Delineating Urban Housing Submarkets with Fuzzy Clustering. *Proceedings of the 15th ACM International Symposium on Geographic Information Systems*, 176-183. Retrieved from [https://gis.depaul.edu/shwang/research/ACMGIS07\\_Paper176\\_HwangThill\\_20070926.pdf](https://gis.depaul.edu/shwang/research/ACMGIS07_Paper176_HwangThill_20070926.pdf)
- [3] L. Zhang and P. L. Jin. (2023). Integrating Clustering and Regression for Spatial Segmentation in Housing Markets. *arXiv Preprint*. Retrieved from <https://arxiv.org/html/2405.08398v2>
- [4] G. S. Sirmans, D. A. Macpherson, and E. N. Zietz. (2005). The Composition of Hedonic Pricing Models. *Journal of Real Estate Literature*, 13(1), 3-43.



# A Code

Listing 1: Python Script for Analysis

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 import folium
6 from folium.plugins import MarkerCluster
7 from sklearn.cluster import KMeans
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.cluster import KMeans
10 import matplotlib.pyplot as plt
11 import plotly.express as px
12
13 # Load the dataset
14 df = pd.read_csv("NY_House_Dataset.csv")
15
16 # === Step 1: Remove Outliers in PRICE ===
17 # Remove properties above the 99th percentile in price
18 upper_price_limit = df["PRICE"].quantile(0.99)
19 df = df[df["PRICE"] <= upper_price_limit]
20
21 # Visualize PRICE distribution after removing outliers
22 plt.figure(figsize=(8, 4))
23 df["PRICE"].plot(kind="box")
24 plt.title("Price_Distribution_After_Outlier_Removal")
25 plt.savefig('Price_Distribution_After_Outlier_Removal.png')
26 plt.show()
27
28 # === Step 2: Remove Invalid or Zero Square Footage ===
29 df = df[df["PROPERTYSQFT"] > 0]
30
31 # === Step 3: Handle Bedrooms (BEDS) and Bathrooms (BATH) ===
32 # Round BEDS and BATH to nearest integer
33 df['BEDS'] = df['BEDS'].round(0).astype(int)
34 df['BATH'] = df['BATH'].round(0).astype(int)
35
36 # Remove unrealistic BEDS and BATH values
37 df = df[(df['BEDS'] >= 1) & (df['BEDS'] <= 15)]
38 df = df[(df['BATH'] >= 1) & (df['BATH'] <= 15)]
39
40 # === Step 4: Handle Invalid or Missing Street Names ===
41 # Identify rows with numeric LONG_NAME values and replace
   them with a placeholder
```

```

42 invalid_street = df['LONG_NAME'].str.isnumeric()
43 df.loc[invalid_street, 'LONG_NAME'] = 'unknown'
44
45 # Replace empty or null strings in STREET_NAME with '
    unknown'
46 df['STREET_NAME'] = df['STREET_NAME'].fillna('unknown').str
    .strip().str.lower()
47 df['LONG_NAME'] = df['LONG_NAME'].fillna('unknown').str
    .strip().str.lower()
48
49 # === Step 5: Remove Duplicates ===
50 df = df.drop_duplicates(subset=["ADDRESS", "PRICE", "TYPE"
    ])
51
52 # === Step 6: Filter Properties by Geographic Boundaries (
    NYC) ===
53 nyc_lat_min, nyc_lat_max = 40.4, 41.0
54 nyc_lon_min, nyc_lon_max = -74.3, -73.6
55 df = df[(df["LATITUDE"].between(nyc_lat_min, nyc_lat_max))
    &
56         (df["LONGITUDE"].between(nyc_lon_min, nyc_lon_max))
    ]
57
58 # === Step 7: Standardize Text Columns ===
59 text_cols = ["BROKERTITLE", "TYPE", "ADDRESS", "STATE", "
    MAIN_ADDRESS",
60              "ADMINISTRATIVE_AREA_LEVEL_2", "LOCALITY", "
    SUBLOCALITY",
61              "STREET_NAME", "LONG_NAME", "FORMATTED_ADDRESS
    "]
62 for col in text_cols:
63     df[col] = df[col].astype(str).str.strip().str.lower()
64
65 # === Step 8: Add Derived Features ===
66 df["PRICE_PER_SQFT"] = df["PRICE"] / df["PROPERTYSQFT"]
67 df = df[df["PRICE_PER_SQFT"] > 0] # Ensure no invalid
    values
68 df["PRICE_PER_SQFT"] = df["PRICE_PER_SQFT"].round(2)
69 df["LOG_PRICE"] = np.log(df["PRICE"]).round(2)
70 df["LOG_PRICE_PER_SQFT"] = np.log(df["PRICE_PER_SQFT"]).
    round(2)
71
72 # === Step 9: Save the Cleaned Dataset ===
73 # df.to_csv("Cleaned_NY_House_Dataset.csv", index=False)
74 # print(f"Cleaned dataset saved with {df.shape[0]} rows and
    {df.shape[1]} columns.")
75

```

```

76 # Print a summary of the cleaned data
77 # print(f"Cleaned data saved with {df.shape[0]} rows and {
    df.shape[1]} columns.")
78 # .info()
79 # print(df.describe())
80
81
82 # Analyze the data
83 locality_price_stats = df.groupby('SUBLOCALITY')['PRICE'].
    agg(['mean', 'median', 'count', 'std']).reset_index()
84 locality_price_stats = locality_price_stats.sort_values(by=
    'mean', ascending=False)
85 # print("Top 10 most expensive sub-localities",
    locality_price_stats.head(5)) # Top 10 most expensive
    sub-localities
86
87
88 # Prepare data for clustering
89 clustering_features = df[['LATITUDE', 'LONGITUDE', 'PRICE'
    ]].copy()
90 scaler = StandardScaler()
91 clustering_features_scaled = scaler.fit_transform(
    clustering_features)
92
93 # Data preparation
94 clustering_features = df[['LATITUDE', 'LONGITUDE', 'PRICE'
    ]]
95 scaler = StandardScaler()
96 clustering_features_scaled = scaler.fit_transform(
    clustering_features)
97
98 # Elbow method
99 wcss = []
100 for k in range(1, 15):
101     kmeans = KMeans(n_clusters=k, random_state=42)
102     kmeans.fit(clustering_features_scaled)
103     wcss.append(kmeans.inertia_)
104
105 # Apply K-Means
106 kmeans = KMeans(n_clusters=6, random_state=42)
107 df['CLUSTER'] = kmeans.fit_predict(
    clustering_features_scaled)
108
109 threshold = 50 # Example: Remove square footage values
    with more than 50 occurrences
110
111 # Identify frequent square footage values

```

```

112 frequent_sqrft_values = df['PROPERTYSQFT'].value_counts()
113 frequent_sqrft_values = frequent_sqrft_values[
    frequent_sqrft_values > threshold].index
114
115 # Remove rows with these square footage values
116 df = df[~df['PROPERTYSQFT'].isin(frequent_sqrft_values)]
117 df.to_csv("NY_House_Dataset_with_clusters.csv", index=False)
118
119 cluster_summary = df.groupby('CLUSTER').agg({
120     'PRICE': ['mean', 'median'],
121     'PRICE_PER_SQFT': ['mean', 'median'],
122     'BEDS': 'mean',
123     'BATH': 'mean'
124 })
125
126 # Expand display options
127 pd.set_option('display.max_rows', 100) # Set to a large
    value to show more rows
128 pd.set_option('display.max_columns', 100) # Set to a large
    value to show more columns
129 borough_cluster_summary = df.groupby(['CLUSTER', '
    SUBLOCALITY']).size().unstack(fill_value=0)
130
131
132 # distribution of PRICE
133 sns.histplot(df["PRICE"], kde=True)
134 plt.title("Distribution of PRICE")
135 plt.show()
136 # distribution of PRICE_PER_SQFT
137 sns.histplot(df["PRICE_PER_SQFT"], kde=True)
138 plt.title("Distribution of PRICE_PER_SQFT")
139 plt.show()
140 # distribution of LOG_PRICE
141 df["LOG_PRICE"] = np.log1p(df["PRICE"])
142 sns.histplot(df["LOG_PRICE"], kde=True)
143 plt.title("Distribution of LOG_PRICE")
144 plt.show()
145
146 # distribution of LOG_PRICE_PER_SQFT
147 corr = df.corr(numeric_only=True)
148 sns.heatmap(corr, annot=True, fmt=".2f", cmap="coolwarm")
149 plt.title("Correlation Matrix")
150 plt.show()
151
152 # Plot the elbow curve
153 plt.figure(figsize=(8, 6))

```

```

154 plt.plot(range(1, 15), wcss, marker='o')
155 plt.title("Elbow Method to Determine Optimal k")
156 plt.xlabel("Number of Clusters (k)")
157 plt.ylabel("WCSS")
158 plt.show()
159
160 # Visualize Clusters
161 plt.figure(figsize=(10, 6))
162 sns.scatterplot(data=df, x='LONGITUDE', y='LATITUDE', hue='
    CLUSTER', palette='tab10', alpha=0.7)
163 plt.title("Clusters of Sub-Localities")
164 plt.show()
165
166
167 # Filter sub-localities with a significant number of
    listings to avoid clutter
168 sublocality_filter = locality_price_stats[
    locality_price_stats['count'] > 10]['SUBLOCALITY']
169 filtered_df = df[df['SUBLOCALITY'].isin(sublocality_filter)
    ]
170
171 plt.figure(figsize=(12, 10))
172 sns.boxplot(data=filtered_df, x='SUBLOCALITY', y='PRICE',
    showfliers=False)
173 plt.xticks(rotation=90)
174 plt.title("Distribution of Prices by Sub-Locality")
175 plt.show()
176
177 plt.figure(figsize=(12, 10))
178 sns.boxplot(data=filtered_df, x='TYPE', y='PRICE',
    showfliers=False)
179 plt.xticks(rotation=90)
180 plt.title("Distribution of Prices by TYPR")
181 plt.show()
182
183 # Create a map
184
185
186 # df['CLUSTER'] = df['CLUSTER'].astype(str)
187 # Define colors for each cluster
188 cluster_colors = {
189     0: 'blue',
190     1: 'orange',
191     2: 'green',
192     3: 'red',
193     4: 'purple',
194     5: 'brown'

```

```

195 }
196
197 # Create the map
198 map = folium.Map(location=[40.7128, -74.0060], zoom_start
199                  =12)
200
201 # Create a MarkerCluster for each cluster
202 for cluster_label, color in cluster_colors.items():
203     # Filter the dataset for this cluster
204     cluster_data = df[df['CLUSTER'] == cluster_label]
205
206     # Create a MarkerCluster for this cluster
207     cluster = MarkerCluster(icon_create_function=f'''
208         function(cluster) {{
209             return L.divIcon({{
210                 html: '<div style="background-color:{color}
211                     }; color:white; border-radius:50%;
212                     padding:5px;">' + cluster.getChildCount
213                     () + '</div>',
214                 className: 'marker-cluster',
215                 iconSize: [30, 30]
216             }});
217         }}
218     ''').add_to(map)
219
220 # Add CircleMarkers for this cluster
221 for i, row in cluster_data.iterrows():
222     folium.CircleMarker(
223         location=(row['LATITUDE'], row['LONGITUDE']),
224         radius=5,
225         color=color,
226         fill=True,
227         fill_opacity=0.7,
228         tooltip=f"Price: ⬇️${row['PRICE']:, .0f}, ⬇️Cluster:
229                 ⬇️{cluster_label}"
230     ).add_to(cluster)
231
232 # Save the map as an interactive HTML file
233 map.save("NY_Housing_Map_Clusters.html")
234 # print("Map saved as 'NY_Housing_Map_Clusters.html'")
235
236 print('cluster_summary: '
237       , cluster_summary)
238 # print('borough_cluster_summary', borough_cluster_summary)

```

```

237 from scipy.stats import f_oneway
238
239 # Test if PRICE differs significantly across clusters
240 anova_result = f_oneway(
241     *[df[df['CLUSTER'] == cluster]['PRICE'] for cluster in
242       df['CLUSTER'].unique()]
243 )
244 print(f"ANOVA_result_for_PRICE_across_clusters:{
245     anova_result}")
246
247 from sklearn.cluster import KMeans
248 from sklearn.metrics import silhouette_score
249
250 # Select features for clustering
251 X = df[['LATITUDE', 'LONGITUDE', 'PRICE']]
252
253 # Drop rows with missing values to ensure consistency
254 X = X.dropna()
255
256 # Fit KMeans on the cleaned dataset
257 kmeans = KMeans(n_clusters=5, random_state=42).fit(X)
258
259 # Ensure the same dataset is used for silhouette score
260 score = silhouette_score(X, kmeans.labels_)
261 print(f'Silhouette_Score:{score}')
262 print(f"Shape_of_X:{X.shape}")
263 print(f"Length_of_labels:{len(kmeans.labels_)}")

```